Deep Learning

Ian Goodfellow
Yoshua Bengio
Aaron Courville

Contents

Website vii

Acknowledgments viii

Notation xi

1 Introduction 1
- 1.1 Who Should Read This Book? . . . 8
- 1.2 Historical Trends in Deep Learning . . . 11

I Applied Math and Machine Learning Basics 29

2 Linear Algebra 31
- 2.1 Scalars, Vectors, Matrices and Tensors . . . 31
- 2.2 Multiplying Matrices and Vectors . . . 34
- 2.3 Identity and Inverse Matrices . . . 36
- 2.4 Linear Dependence and Span . . . 37
- 2.5 Norms . . . 39
- 2.6 Special Kinds of Matrices and Vectors . . . 40
- 2.7 Eigendecomposition . . . 42
- 2.8 Singular Value Decomposition . . . 44
- 2.9 The Moore-Penrose Pseudoinverse . . . 45
- 2.10 The Trace Operator . . . 46
- 2.11 The Determinant . . . 47
- 2.12 Example: Principal Components Analysis . . . 48

3 Probability and Information Theory 53
- 3.1 Why Probability? . . . 54

3.2 Random Variables . . . 56
3.3 Probability Distributions . . . 56
3.4 Marginal Probability . . . 58
3.5 Conditional Probability . . . 59
3.6 The Chain Rule of Conditional Probabilities . . . 59
3.7 Independence and Conditional Independence . . . 60
3.8 Expectation, Variance and Covariance . . . 60
3.9 Common Probability Distributions . . . 62
3.10 Useful Properties of Common Functions . . . 67
3.11 Bayes' Rule . . . 70
3.12 Technical Details of Continuous Variables . . . 71
3.13 Information Theory . . . 73
3.14 Structured Probabilistic Models . . . 75

4 Numerical Computation 80
4.1 Overflow and Underflow . . . 80
4.2 Poor Conditioning . . . 82
4.3 Gradient-Based Optimization . . . 82
4.4 Constrained Optimization . . . 93
4.5 Example: Linear Least Squares . . . 96

5 Machine Learning Basics 98
5.1 Learning Algorithms . . . 99
5.2 Capacity, Overfitting and Underfitting . . . 110
5.3 Hyperparameters and Validation Sets . . . 120
5.4 Estimators, Bias and Variance . . . 122
5.5 Maximum Likelihood Estimation . . . 131
5.6 Bayesian Statistics . . . 135
5.7 Supervised Learning Algorithms . . . 140
5.8 Unsupervised Learning Algorithms . . . 146
5.9 Stochastic Gradient Descent . . . 151
5.10 Building a Machine Learning Algorithm . . . 153
5.11 Challenges Motivating Deep Learning . . . 155

II Deep Networks: Modern Practices 166

6 Deep Feedforward Networks 168
6.1 Example: Learning XOR . . . 171
6.2 Gradient-Based Learning . . . 177

6.3 Hidden Units . . . 191
6.4 Architecture Design . . . 197
6.5 Back-Propagation and Other Differentiation Algorithms . . . 204
6.6 Historical Notes . . . 224

7 Regularization for Deep Learning 228
7.1 Parameter Norm Penalties . . . 230
7.2 Norm Penalties as Constrained Optimization . . . 237
7.3 Regularization and Under-Constrained Problems . . . 239
7.4 Dataset Augmentation . . . 240
7.5 Noise Robustness . . . 242
7.6 Semi-Supervised Learning . . . 243
7.7 Multi-Task Learning . . . 244
7.8 Early Stopping . . . 246
7.9 Parameter Tying and Parameter Sharing . . . 253
7.10 Sparse Representations . . . 254
7.11 Bagging and Other Ensemble Methods . . . 256
7.12 Dropout . . . 258
7.13 Adversarial Training . . . 268
7.14 Tangent Distance, Tangent Prop, and Manifold Tangent Classifier 270

8 Optimization for Training Deep Models 274
8.1 How Learning Differs from Pure Optimization . . . 275
8.2 Challenges in Neural Network Optimization . . . 282
8.3 Basic Algorithms . . . 294
8.4 Parameter Initialization Strategies . . . 301
8.5 Algorithms with Adaptive Learning Rates . . . 306
8.6 Approximate Second-Order Methods . . . 310
8.7 Optimization Strategies and Meta-Algorithms . . . 317

9 Convolutional Networks 330
9.1 The Convolution Operation . . . 331
9.2 Motivation . . . 335
9.3 Pooling . . . 339
9.4 Convolution and Pooling as an Infinitely Strong Prior . . . 345
9.5 Variants of the Basic Convolution Function . . . 347
9.6 Structured Outputs . . . 358
9.7 Data Types . . . 360
9.8 Efficient Convolution Algorithms . . . 362
9.9 Random or Unsupervised Features . . . 363

9.10 The Neuroscientific Basis for Convolutional Networks 364
9.11 Convolutional Networks and the History of Deep Learning 371

10 Sequence Modeling: Recurrent and Recursive Nets 373
10.1 Unfolding Computational Graphs 375
10.2 Recurrent Neural Networks 378
10.3 Bidirectional RNNs 394
10.4 Encoder-Decoder Sequence-to-Sequence Architectures 396
10.5 Deep Recurrent Networks 398
10.6 Recursive Neural Networks 400
10.7 The Challenge of Long-Term Dependencies 401
10.8 Echo State Networks 404
10.9 Leaky Units and Other Strategies for Multiple Time Scales 406
10.10 The Long Short-Term Memory and Other Gated RNNs 408
10.11 Optimization for Long-Term Dependencies 413
10.12 Explicit Memory 416

11 Practical Methodology 421
11.1 Performance Metrics 422
11.2 Default Baseline Models 425
11.3 Determining Whether to Gather More Data 426
11.4 Selecting Hyperparameters 427
11.5 Debugging Strategies 436
11.6 Example: Multi-Digit Number Recognition 440

12 Applications 443
12.1 Large-Scale Deep Learning 443
12.2 Computer Vision 452
12.3 Speech Recognition 458
12.4 Natural Language Processing 461
12.5 Other Applications 478

III Deep Learning Research 486

13 Linear Factor Models 489
13.1 Probabilistic PCA and Factor Analysis 490
13.2 Independent Component Analysis (ICA) 491
13.3 Slow Feature Analysis 493
13.4 Sparse Coding 496

13.5 Manifold Interpretation of PCA 499

14 Autoencoders **502**
14.1 Undercomplete Autoencoders 503
14.2 Regularized Autoencoders 504
14.3 Representational Power, Layer Size and Depth 508
14.4 Stochastic Encoders and Decoders 509
14.5 Denoising Autoencoders 510
14.6 Learning Manifolds with Autoencoders 515
14.7 Contractive Autoencoders 521
14.8 Predictive Sparse Decomposition 523
14.9 Applications of Autoencoders 524

15 Representation Learning **526**
15.1 Greedy Layer-Wise Unsupervised Pretraining 528
15.2 Transfer Learning and Domain Adaptation 536
15.3 Semi-Supervised Disentangling of Causal Factors 541
15.4 Distributed Representation 546
15.5 Exponential Gains from Depth 553
15.6 Providing Clues to Discover Underlying Causes 554

16 Structured Probabilistic Models for Deep Learning **558**
16.1 The Challenge of Unstructured Modeling 559
16.2 Using Graphs to Describe Model Structure 563
16.3 Sampling from Graphical Models 580
16.4 Advantages of Structured Modeling 582
16.5 Learning about Dependencies 582
16.6 Inference and Approximate Inference 584
16.7 The Deep Learning Approach to Structured Probabilistic Models 585

17 Monte Carlo Methods **590**
17.1 Sampling and Monte Carlo Methods 590
17.2 Importance Sampling 592
17.3 Markov Chain Monte Carlo Methods 595
17.4 Gibbs Sampling 599
17.5 The Challenge of Mixing between Separated Modes 599

18 Confronting the Partition Function **605**
18.1 The Log-Likelihood Gradient 606
18.2 Stochastic Maximum Likelihood and Contrastive Divergence . . . 607

18.3 Pseudolikelihood . . . 615
18.4 Score Matching and Ratio Matching . . . 617
18.5 Denoising Score Matching . . . 619
18.6 Noise-Contrastive Estimation . . . 620
18.7 Estimating the Partition Function . . . 623

19 Approximate Inference **631**
19.1 Inference as Optimization . . . 633
19.2 Expectation Maximization . . . 634
19.3 MAP Inference and Sparse Coding . . . 635
19.4 Variational Inference and Learning . . . 638
19.5 Learned Approximate Inference . . . 651

20 Deep Generative Models **654**
20.1 Boltzmann Machines . . . 654
20.2 Restricted Boltzmann Machines . . . 656
20.3 Deep Belief Networks . . . 660
20.4 Deep Boltzmann Machines . . . 663
20.5 Boltzmann Machines for Real-Valued Data . . . 676
20.6 Convolutional Boltzmann Machines . . . 683
20.7 Boltzmann Machines for Structured or Sequential Outputs . . . 685
20.8 Other Boltzmann Machines . . . 686
20.9 Back-Propagation through Random Operations . . . 687
20.10 Directed Generative Nets . . . 692
20.11 Drawing Samples from Autoencoders . . . 711
20.12 Generative Stochastic Networks . . . 714
20.13 Other Generation Schemes . . . 716
20.14 Evaluating Generative Models . . . 717
20.15 Conclusion . . . 720

Bibliography **721**

Index **777**

Website

www.deeplearningbook.org

This book is accompanied by the above website. The website provides a variety of supplementary material, including exercises, lecture slides, corrections of mistakes, and other resources that should be useful to both readers and instructors.

Acknowledgments

This book would not have been possible without the contributions of many people.

We would like to thank those who commented on our proposal for the book and helped plan its contents and organization: Guillaume Alain, Kyunghyun Cho, Çağlar Gülçehre, David Krueger, Hugo Larochelle, Razvan Pascanu and Thomas Rohée.

We would like to thank the people who offered feedback on the content of the book itself. Some offered feedback on many chapters: Martín Abadi, Guillaume Alain, Ion Androutsopoulos, Fred Bertsch, Olexa Bilaniuk, Ufuk Can Biçici, Matko Bošnjak, John Boersma, Greg Brockman, Alexandre de Brébisson, Pierre Luc Carrier, Sarath Chandar, Pawel Chilinski, Mark Daoust, Oleg Dashevskii, Laurent Dinh, Stephan Dreseitl, Jim Fan, Miao Fan, Meire Fortunato, Frédéric Francis, Nando de Freitas, Çağlar Gülçehre, Jurgen Van Gael, Javier Alonso García, Jonathan Hunt, Gopi Jeyaram, Chingiz Kabytayev, Lukasz Kaiser, Varun Kanade, Asifullah Khan, Akiel Khan, John King, Diederik P. Kingma, Yann LeCun, Rudolf Mathey, Matías Mattamala, Abhinav Maurya, Kevin Murphy, Oleg Mürk, Roman Novak, Augustus Q. Odena, Simon Pavlik, Karl Pichotta, Eddie Pierce, Kari Pulli, Roussel Rahman, Tapani Raiko, Anurag Ranjan, Johannes Roith, Mihaela Rosca, Halis Sak, César Salgado, Grigory Sapunov, Yoshinori Sasaki, Mike Schuster, Julian Serban, Nir Shabat, Ken Shirriff, Andre Simpelo, Scott Stanley, David Sussillo, Ilya Sutskever, Carles Gelada Sáez, Graham Taylor, Valentin Tolmer, Massimiliano Tomassoli, An Tran, Shubhendu Trivedi, Alexey Umnov, Vincent Vanhoucke, Marco Visentini-Scarzanella, Martin Vita, David Warde-Farley, Dustin Webb, Kelvin Xu, Wei Xue, Ke Yang, Li Yao, Zygmunt Zając and Ozan Çağlayan.

We would also like to thank those who provided us with useful feedback on individual chapters:

- Notation: Zhang Yuanhang.
- Chapter 1, Introduction: Yusuf Akgul, Sebastien Bratieres, Samira Ebrahimi,

Charlie Gorichanaz, Brendan Loudermilk, Eric Morris, Cosmin Pârvulescu and Alfredo Solano.

- Chapter 2, Linear Algebra: Amjad Almahairi, Nikola Banić, Kevin Bennett, Philippe Castonguay, Oscar Chang, Eric Fosler-Lussier, Andrey Khalyavin, Sergey Oreshkov, István Petrás, Dennis Prangle, Thomas Rohée, Gitanjali Gulve Sehgal, Colby Toland, Alessandro Vitale and Bob Welland.

- Chapter 3, Probability and Information Theory: John Philip Anderson, Kai Arulkumaran, Vincent Dumoulin, Rui Fa, Stephan Gouws, Artem Oboturov, Antti Rasmus, Alexey Surkov and Volker Tresp.

- Chapter 4, Numerical Computation: Tran Lam AnIan Fischer and Hu Yuhuang.

- Chapter 5, Machine Learning Basics: Dzmitry Bahdanau, Justin Domingue, Nikhil Garg, Makoto Otsuka, Bob Pepin, Philip Popien, Emmanuel Rayner, Peter Shepard, Kee-Bong Song, Zheng Sun and Andy Wu.

- Chapter 6, Deep Feedforward Networks: Uriel Berdugo, Fabrizio Bottarel, Elizabeth Burl, Ishan Durugkar, Jeff Hlywa, Jong Wook Kim, David Krueger and Aditya Kumar Praharaj.

- Chapter 7, Regularization for Deep Learning: Morten Kolbæk, Kshitij Lauria, Inkyu Lee, Sunil Mohan, Hai Phong Phan and Joshua Salisbury.

- Chapter 8, Optimization for Training Deep Models: Marcel Ackermann, Peter Armitage, Rowel Atienza, Andrew Brock, Tegan Maharaj, James Martens, Kashif Rasul, Klaus Strobl and Nicholas Turner.

- Chapter 9, Convolutional Networks: Martín Arjovsky, Eugene Brevdo, Konstantin Divilov, Eric Jensen, Mehdi Mirza, Alex Paino, Marjorie Sayer, Ryan Stout and Wentao Wu.

- Chapter 10, Sequence Modeling: Recurrent and Recursive Nets: Gökçen Eraslan, Steven Hickson, Razvan Pascanu, Lorenzo von Ritter, Rui Rodrigues, Dmitriy Serdyuk, Dongyu Shi and Kaiyu Yang.

- Chapter 11, Practical Methodology: Daniel Beckstein.

- Chapter 12, Applications: George Dahl, Vladimir Nekrasov and Ribana Roscher.

- Chapter 13, Linear Factor Models: Jayanth Koushik.